A COPILOT PRICING THOUGHT-EXPERIMENT · REVISED

If GitHub dropped PRU billing for tokens, who actually wins?

Copilot meters most models in *premium request units* — fixed-price chats with a per-model multiplier. Direct API providers meter in tokens. The two systems disagree by an order of magnitude in opposite directions, and which way they disagree depends entirely on how you use Copilot. The first version of this analysis only modeled chat. The corrected version below models **chat, agent mode, Copilot CLI, sub-agents and the cloud agent** — and the verdict reverses for most users.

HEAVIEST PRU MARKUP (CHAT ONLY)

4.3×

Opus 4.7 PRU vs simple-chat tokens

HEAVIEST PRU SUBSIDY (CLOUD AGENT)

33333×

Same model, large cloud-agent task

WORKFLOW MODES COVERED

7

From simple chat to large cloud-agent runs

ANTHROPIC-PUBLISHED ANCHOR

15×

Multi-agent tokens vs. plain chat

1 The question, revised

On a paid Copilot plan today, a chat with **GPT-4o costs $0**, a chat with **Claude Opus 4.7 costs $0.30**, and the actual provider token cost of those two conversations is roughly **$0.03** and **$0.07**. PRU and tokens disagree by an order of magnitude in opposite directions. So — what would actually move if GitHub flipped the switch?

> **Two-part answer.**
>
> **For chat:**
>
> ultra-premium gets ~4× cheaper, the free tier disappears, and Gemini wins the mid-tier.
>
> **For agentic workflows:**
>
> the math reverses completely — PRU's "tool calls don't count" rule is worth $300–$2,000/month for agent-heavy users, and unwinding it would dwarf every other change.

2 The five PRU tiers, and the token costs that explain them

Copilot collapses 18 models into **five pricing tiers**. Each tier's multiplier is a deliberate choice — sometimes a fair reflection of the model's underlying token cost, sometimes a strategic subsidy or markup. Here's what each tier actually buys you and why GitHub almost certainly priced it that way.

Included 0×

PRU charge: **$0.00**/chat vs token cost: **$0.005–$0.030**/chat

`GPT-5 mini` `GPT-4.1` `GPT-4o`

Zero marginal cost per chat — bundled into the seat fee, which more than covers negotiated provider rates at typical usage. The $0.005–$0.03/chat figure is provider list price, not GitHub's actual cost-to-serve.

Zero-marginal-cost bundle. User-perceived subsidy at list price; on GitHub's books the seat fee covers negotiated inference costs.

Budget Premium 0.25×

PRU charge: **$0.01**/chat vs token cost: **$0.012–$0.014**/chat

`Claude Haiku 4.5` `GPT-5.4 mini` `Grok Code Fast 1`

Roughly fair pricing. The 0.25× multiplier closely tracks actual token cost for short chats.

Within 30% of token cost. Honest pricing.

Standard Premium 1×

PRU charge: **$0.04**/chat vs token cost: **$0.025–$0.042**/chat

`GPT-5.2` `Claude Sonnet 4.5` `Gemini 3.1 Pro`

Slight chat-mode overcharge for cheaper-token models (Gemini at $0.025), roughly fair for Anthropic Sonnet at $3/$15. Becomes a subsidy in agent modes.

Chat: PRU charges 5–60% more than tokens. Agent mode: PRU is 5–10× cheaper than tokens.

High Premium 1.25×

PRU charge: **$0.05**/chat vs token cost: **$0.040–$0.070**/chat

`Claude Opus 4.5` `Claude Opus 4.6` `GPT-5.4`

Significant chat-mode subsidy: Opus actually costs $0.07/chat at list price, charged at $0.05. The cloud-agent subsidy reaches 200× for large tasks.

Always favors the user. 1.4× chat subsidy, 10–200× agent subsidy.

Ultra Premium 7.5×

PRU charge: **$0.30**/chat vs token cost: **$0.070–$0.340**/chat

`Claude Opus 4.7` `GPT-5.5`

Designed to fund the rest of the system. Looks like a 4× markup against simple chats but is near-fair against workspace chats and a subsidy against any agent workflow.

Chat: PRU charges 1.4–4× tokens. Cloud agent: PRU is 11–33× cheaper than tokens.

> **Read this as a system, not a list.**
> The 7.5× ultra-premium tier *funds* the 0× included tier. The 1.25× high-premium tier is consistently subsidized at chat sizes *and* at agent sizes — a deliberate "use Opus, don't be afraid" signal. Standard Premium (1×) is the only tier that's roughly cost-neutral against tokens.

3 Every model placed on the multiplier ladder

Eighteen models, five tiers. Hover any bar to see the multiplier, the per-chat PRU charge, and the underlying token cost for that model.

Copilot premium-request multiplier by model

Higher = more PRU per chat. Bars colored by tier. Multiplier↔list-price-token-cost correlation across these models is **r ≈ 0.80** — real but loose.

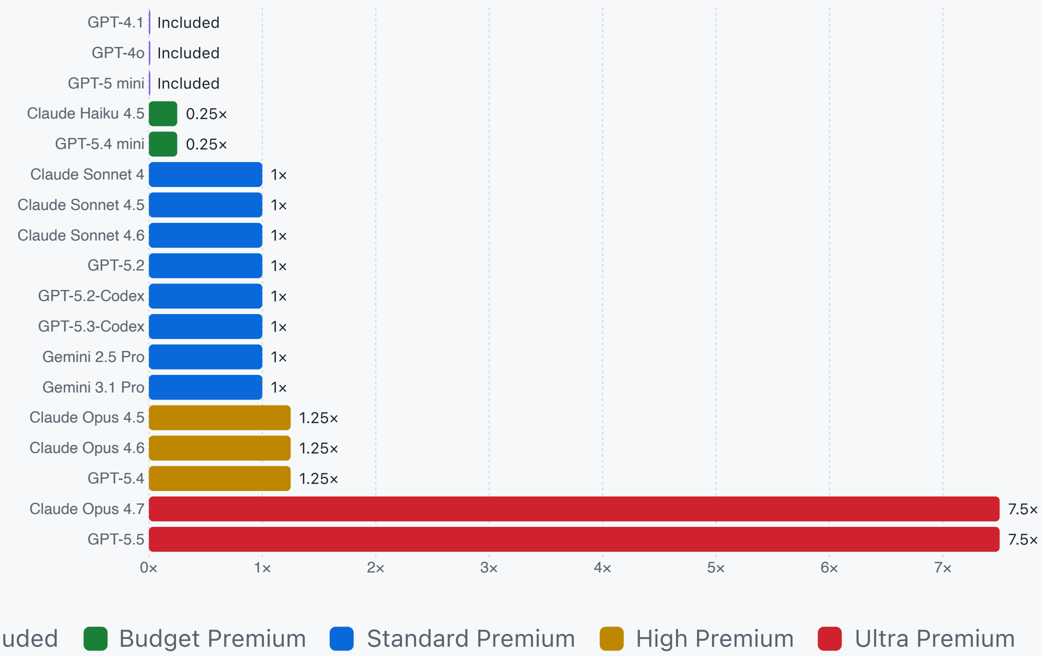

Source: docs.github.com — Copilot supported models

4 The hidden subsidy — and why the chat picture is wrong

Subtracting the actual provider-token cost of a typical chat from the PRU charge reveals who is over- and under-charged. **The first toggle below is the original analysis (4K-in / 2K-out simple chat). The second toggle re-runs the same comparison at agent-mode token volumes** (60K in / 15K out). Watch the ultra-premium bars flip sign.

Read this as a *user-perceived* subsidy, not a P&L subsidy.

Every "subsidy" figure on this page benchmarks the PRU charge against *provider list price* (Anthropic, OpenAI, Google public API rates). GitHub's actual cost-to-serve is materially lower because of negotiated commit rates, the Microsoft↔OpenAI revenue-share, Azure-hosted inference, and provider-side prompt caching / batch discounts. The *relative ordering* of subsidies across tiers is preserved, but the *absolute* GitHub→user transfer is smaller than the bars below imply — and on paid plans the included tier's "zero marginal cost" is bundled into the seat fee, not paid for by GitHub out of pocket.

Simple chat (4K + 2K) | Agent mode (60K + 15K)

PRU price minus provider list-price token cost · per task

Right (red) = PRU charges more than list price · Left (green) = PRU charges less than list price (user-perceived subsidy). GitHub's negotiated cost-to-serve is lower than list, so absolute green-side magnitudes overstate the P&L cost to GitHub.

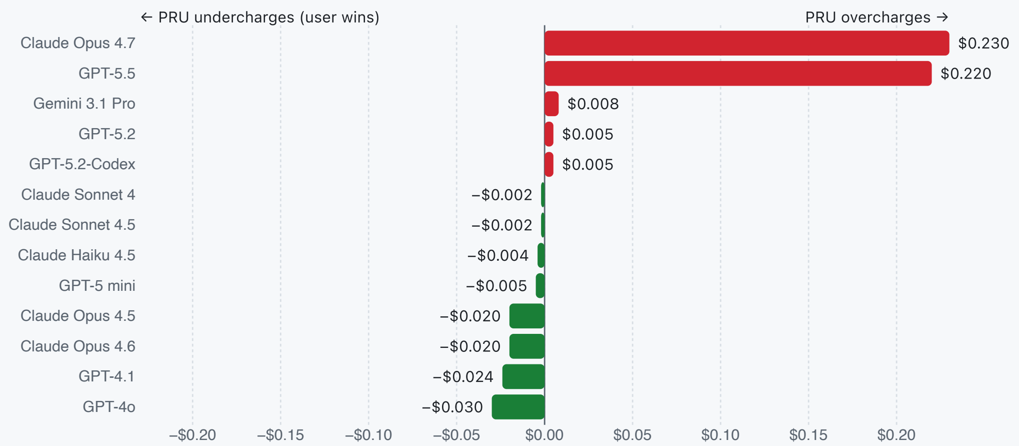

The "ultra-premium overcharges users" finding only holds for short chats.

The moment a user prompt triggers any kind of agentic loop — workspace context, tool calls, tests, sub-agents — the high and ultra tiers flip from the red side to the green side. PRU subsidizes *users* of those tiers in the modes those tiers were designed for. Note also that the 1.25×→7.5× jump is a

6× price step against a 1× step in list-price token cost

— that gap reflects reasoning/adaptive-thinking output volume, promotional launch pricing (Opus 4.7's 7.5× rate expires April 30, 2026), and deliberate cross-subsidy of the free and agentic tiers, not raw input/output rates.

5 Why "chat" is the wrong unit

Copilot in 2026 is not one product — it's at least **seven workflow modes** spanning nearly three orders of magnitude in token consumption. The 4K-in / 2K-out "typical chat" figure that the original PRU vs. token comparison rested on describes only the leftmost column of this picture.

Token volume by workflow mode (log scale)

One bar = one full task or session. The same prompt costs the same in PRU; under tokens it scales with the bar.

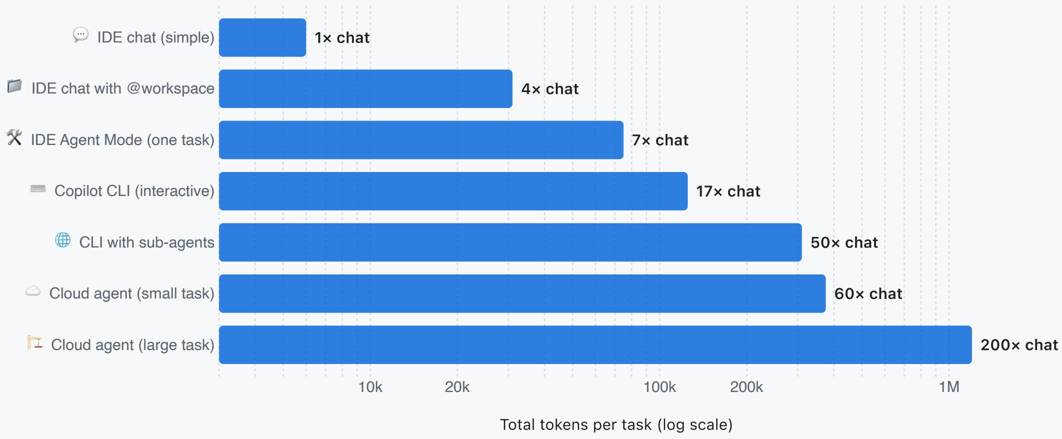

Anchored to Anthropic's published 4× / 15× chat-vs-agent multipliers and SWE-bench bash-only $/instance measured costs.

💬 **IDE chat (simple)**
Stand-alone Q&A, single file open. The original 'baseline chat' the analysis assumed.
6k

📁 **IDE chat with @workspace**
RAG-retrieved chunks dominate input. Gemini-style 'open many files' workflows.
31k

🛠 **IDE Agent Mode (one task)**
5–30 autonomous tool calls per user prompt. Each turn re-sends growing history.
75k

⌨ **Copilot CLI (interactive)**
10–100 turns with plan mode + bash. The terminal-native agent.
125k

🌐 **CLI with sub-agents**
Lead agent + 3–5 parallel sub-agents. Anthropic's measured 15× chat token figure.
310k

☁ **Cloud agent (small task)**
Bug fixes, doc updates. Plan → research → implement → test → PR loop.
375k

🏗 **Cloud agent (large task)**
Refactors, migrations. 50–500 turns on a GitHub Actions runner.
1.2M

The "tool calls don't count" rule.

The Copilot billing docs are explicit: *"only the prompts you send count as premium requests; actions Copilot takes autonomously to complete your task, such as tool calls, do not."* That single sentence is the most valuable subsidy in the entire Copilot pricing system, and it grows non-linearly with workflow autonomy.

6 Cost vs. SWE-bench performance

The headline trade-off: how much does each model cost *per chat*, and how does that map to its **SWE-bench bash-only resolve rate**? Toggle between PRU pricing and token pricing to watch the frontier reshuffle.

PRU pricing (today) | Token pricing (hypothetical)

SWE-bench resolve rate vs. cost per typical chat

Up-and-to-the-left is best. Bubble color = provider.

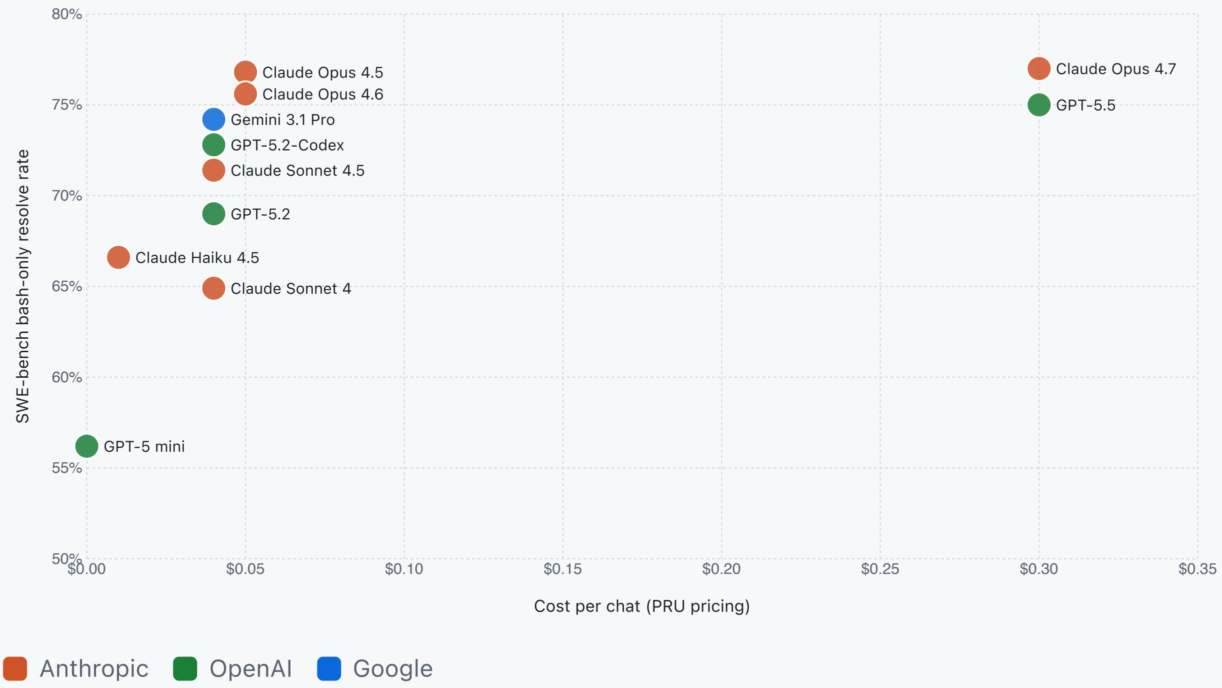

7 Run your own numbers — with workflow mode

Pick a model, a workflow mode, and a monthly volume. The chart compares your bill under PRU billing (where the workflow mode is irrelevant) versus token billing (where it dominates).

MODEL

Claude Opus 4.7 · Ultra Premium · 7.5×

WORKFLOW MODE

IDE chat (simple) (4k in / 2k out)

Stand-alone Q&A, single file open. The original 'baseline chat' the analysis assumed.

TASKS / CHATS PER MONTH (200)

OVERRIDE INPUT TOKENS / TASK (4,000)

OVERRIDE OUTPUT TOKENS / TASK (2,000)

PRU MONTHLY	TOKEN MONTHLY	DELTA (PRU – TOKENS)
$60	$14	+$46

8 Winners and losers under token billing

The corrected impact analysis. Each row models a real user profile under the *actual* workflow mix that profile uses — not a chat-only assumption. Bars show the monthly-cost range under token pricing (purple) versus today's PRU bill (blue).

Monthly Copilot cost: today (PRU) vs. hypothetical (tokens) by user profile

Blue bar = today's PRU charge. Purple range = token-pricing low/high estimate. Right-of-blue = the user pays more under tokens.

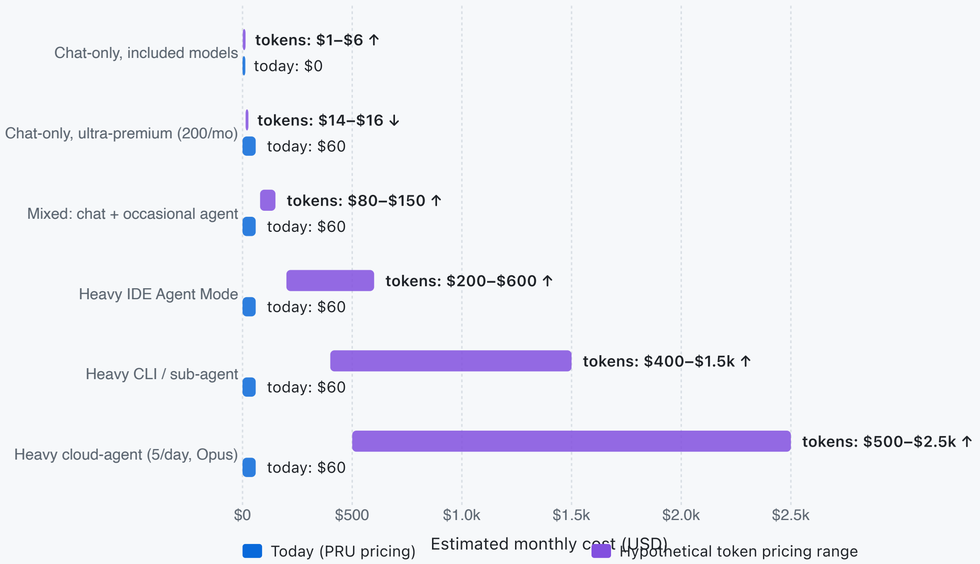

9 The verdict, rewritten

The PRU model is a deliberate cross-subsidy: ultra-premium chat overcharges fund the zero-cost included tier, and PRU's "tool calls don't count" rule funds every agentic workflow on every tier. A token switch would unwind both subsidies — and the second one is much, much larger than the first.

✅ Chat-only ultra-premium users save

Opus 4.7 / GPT-5.5 in plain chat mode drop from $0.30 to ~$0.07–0.08 per turn. A 200-prompt user saves ~$44/month.

⚠️ The free tier disappears

GPT-5 mini, GPT-4.1, GPT-4o all start metering at $0.005–$0.030/chat. Casual users on paid plans pay something for the first time.

⚠️ Agent mode users lose hundreds

One IDE agent task burns 5–10× a chat in tokens. Heavy users pay $200–600/mo extra. Sub-agent and cloud-agent users pay $400–2,500/mo extra.

⚠️ The cloud agent becomes brutal

One large cloud-agent task = $10+ of tokens at provider list price, vs. $0.30 in PRU. Five tasks/day on Opus = ~$1,500/month extra.

🏷️ Cheap models burn more tokens

Gemini 3 Flash at $0.50/MTok solves the same SWE-bench task as Opus at 2.5× the token count. Token pricing rewards iteration efficiency, not just rate.

💾 Caching becomes the budget lever

Anthropic / Gemini cached input is 90% off. A stable 10K-token system prompt becomes effectively free under tokens — PRU never rewarded this.

Bottom line.

Two subsidies hold up Copilot's pricing today. The *ultra-premium overcharge* funds the free tier, and is the one the original analysis surfaced. The *"tool calls don't count" rule* is the much bigger subsidy — it funds every autonomous workflow on every tier, and it's invisible in any per-chat comparison. Token pricing would expose both. Casual chat users would notice; agent-heavy users would feel it like a tax.

Source analysis · SWE-bench · Copilot billing · Anthropic on multi-agent tokens · Cloud agent docs